



**WEST UNIVERSITY OF TIMIȘOARA
FACULTY OF COMPUTER SCIENCE
MASTER STUDY PROGRAM: CYBERSECURITY**

MASTER THESIS

SUPERVISOR:
Dr. Lect. Darius Galis

GRADUATE:
Paul - Andrei Abrudan

**TIMIȘOARA
2026**

WEST UNIVERSITY OF TIMIȘOARA
FACULTY OF COMPUTER SCIENCE
MASTER STUDY PROGRAM: CYBERSECURITY

Automated Security Analysis of the Rust Crate Ecosystem: A Distributed Framework for Vulnerability and Malicious Code Detection

SUPERVISOR:
Dr. Lect. Darius Galis

GRADUATE:
Paul - Andrei Abrudan

TIMIȘOARA
2026

Abstract

The Rust programming language has gained popularity in multiple programming fields due to its memory safety features and performance. Rust's package ecosystem, crates.io, delivers over 150,000 crates to the language's users. Due to the rapid growth of the language and ecosystem, potential security risks are present, including known vulnerabilities in dependencies, a history of supply chain attacks through code injection and typosquatting, secrets being leaked in source code and cases of malicious build scripts present in the optional build.rs file at the root of the packages. This thesis presents a distributed framework for automating security analysis on the Rust package library ecosystem. The system parses metadata from the crates.io database dump, constructing a dependency graph and applying multiple analysis modules: Vulnerability scanning using 'cargo audit', which uses the RustSec Advisory Database, Secret leakage detection using the FOSS tool GitLeaks, heuristic analysis for the 'build.rs' script, detecting network calls, process creation, obfuscation and suspicious TLDs used, the possibility of LLM code scoring using LM Studio's REST API, typosquatting detection using multiple string similarity algorithms and transitive dependency chain risk propagation. Horizontal scaling is possible across multiple nodes, achieved with Redis used as a work queue, with the possibility of upgrading to a faster in-memory database. The results are persisted in a PostgreSQL database and can be visualized through the web dashboard. This thesis aims to make it easier to search for cases of malicious dependencies in the Rust ecosystem, providing tools for security auditing.

Keywords: crates.io, Rust, supply chain security, typosquatting, distributed systems, software security

Statement of Scientific integrity

I declare that this BSc / MSc thesis is an original report and represents my own work, conducted under supervision of [name(s) supervisor(s)]. All research methods, data collection, analysis and interpretation presented in this thesis are accurately reported and all sources have been cited appropriately where I have made use of ideas and/or visuals of others.

If you did not use generative AI tools: I did not use generative AI during the writing and research process.

If you used generative AI tools: I used generative AI during the writing and research process. The details are in the Appendix ??.

Contents

Chapter 1

Introduction

1.1 Context and Motivation

All relatively modern software relies on thirdparty packages, distributed through different centralized registries. The Rust programming language has grown from it's initial release by Mozilla back in 2015, to today where crates.io hosts over 150.000 crates with multiple billions of downloads, being the pillar of the Rust library supply chain.

This growth can be a double-edged sword, while being beneficial for productivity, it introduces an immense attack surface, which remains underexplored as of now. In most languages, including Rust, when a developer adds a dependency to their project, they not only trust the direct dependency but also all its entire dependency tree, which can contain hundreds if not thousands of packages. A single compromised dependency, wether due to an unpatched vulnerability, a typosquat or due to an adversary takeover, can propagate malicious code across a wide part of the ecosystem.

Relatively recent incidents in other package ecosystems, such as npm's attack in 2018 known as 'event-stream', 'ua-parser-js' from 2021 and the 'colors.js' and 'faker.js' sabotages in 2022, demonstrated that supply chain attacks are practical attacks, very effective and can become extremely complex. Rust's memory safety gurantees mitigate some classes of vulnerabilities (such as buffer overflows), they provide no help against attacks such as memory-safe malicious code and secret exfiltration from the environment.

Another attack vector specific to some languages, including Rust, is that the build system allows arbitrary code execution by crates at compile time, using 'build.rs' scripts. The scripts can perform operations such as network communication, process execution and have file system access. This mechanism is essential for some usecases such as native library linking and code generation, however it also is a dangerous attack vector that can infect a machine just compiling a malicious crate, without ever executing the resulted binary.

The motivation comes from the observation that some security tools already exist for specific Rust crate analysis, however no ecosystem-wide comprehensive framework exists, that integrates multiple modules and can apply them at scale.

1.2 Problem Statement

Rust’s crate ecosystem, lacks a scalable analysis framework that is capable of performing module-based scanning, first, the scale of crates.io is massive, with over 150.000+ crates, any ecosystem-wide analysis requires a scalable and distributed architecture, capable of processing crates efficiently. Individual tools address specific parts of security analysis, such as secrets scanning and existing known vulnerabilities. There is no existing solution that combines build script analysis, malicious code detection, typosquatting crate name analysis, vulnerability scanning and secret detection into a complete comprehensive solution. It is also the case that a crate with no direct vulnerabilities can still pose a risk if it exports an API dependent on a vulnerable package, tracking these risks is still an existing challenge. The ‘build.rs’ script allows arbitrary code execution at compile time and it is the case that no study has tracked possibly suspicious build script behaviours across the entire crates.io database. While typosquatting has been studied in other package libraries such as npm and PyPI, the Rust’s ecosystem has seen no studies regarding this. Large Language Models have demonstrated capabilities in understanding code, however the deployment of LLM analysis at scale has not been explored over package libraries.

1.3 Objectives

The objective of this thesis is to design, implement and evaluate a toolkit facilitating comprehensive distributed security analysis over the Rust’s crate ecosystem, the specific objectives are:

Designing a modular toolkit, which unifies multiple modules and analysis techniques into a single pipeline, enabling user friendly security assessment of crates published on crates.io Implementing distributed processing over the entire toolkit, using a ”in-memory database”-based work queue, together with topological ordering of the crates dependencies, we enable horizontal scaling across workers and achieving acceptable throughput on consumer hardware. Develop and integrate multiple analysis modules such as: vulnerability scanning using RustSec’s Database, secret detection in available source code, heuristic ‘build.rs’ analysis for detecting unusual behaviours such as network communication, process execution, obfuscated/encrypted code and suspicious TLD usage, LLM-based malicious code scoring, typosquatting detection using multiple similarity algorithms (evenshtein, Damerau-Levenshtein, Jaro-Winkler, keyboard proximity distance, and prefix), and dependency chain propagation analysis. Conduct an entire ecosystem-wide analysis, either by running the framework on a subset of the crates.io registry, or the entire database, characterizing and doing manual analysis on provided warnings. Develop a web-based dashboard for user friendly exploration and visualization of analysis results. Evaluate the framework’s performance and scalability, measuring throughput, bottlenecks and possible improvements.

1.4 Scope and Delimitations

TODO: Formatting, cleaning, typos, etc

In scope:

- Analysis of multiple crates published on crates.io, imported from the official database dump
- Multiple analysis modules as described in the objective list
- Analysis of crate source code from the git repository
- Having distributed processing capabilities using an in-memory database for work distribution and PostgreSQL for resulting data
- A web dashboard for visualization
- Performance evaluation of analysis and of found problems and exploits

Out of scope:

- Dynamic analysis or runtime behavior
- Analysis of compiled binaries
- Manual verification of all results, due to the framework being designed as a screening tool, requiring human validation
- Real-time monitoring
- Analysis of other package ecosystems
- Formal proofs for detection correctness

1.5 Thesis Structure

This thesis is organized in this way:

Chapter 2 - State of the art: Is a review about existing research and literature on software supply chain security, malicious code detection techniques, package vulnerabilities, typosquatting and build scripts used maliciously. It identifies a research gap related to the previously mentioned objectives.

Chapter 3 - Methodology: Describes the research, data collection and importing, framework design, analysis modules and the evaluation metrics used.

Chapter 4 - System architecture: We detail the high level architecture, covering the data importing, dependency graph creation, the design of the work queue, database schema, the dashboard and deployment infrastructure and optimization scripts.

Chapter 5 - Analysis Modules: We detail the analysis modules used, algorithms, heuristics, fault tolerance and optimization.

Chapter 6 - Implementation Details - We present the technology stack, the CLI app usage, the queue implementation and database implementation

Chapter 7 - Results & Evaluation - This chapter is about the findings resulted by running the framework, vulnerability, leaked secrets, build system statistics, typosquatting possibilities, dependency chain risks and performance benchmarks.

Chapter 8 - Discussion: We interpret the results, compare them to existing literature, discuss limitations and address threats to validity.

Chapter 9 - Conclusion & Future Work: We present a summary of the contribution of this thesis and state possible directions for future research or development

Chapter 2

State of the art

2.1 Software Supply Chain Security

2.1.1 Historical Context and Evolution

Software supply chain security can be considered one of the most critical problems in software engineering, it reshaped how software developers and organizations from the field approach software development, deployment and risk assessment. The concept was introduced by Ken Thompson in the "Reflections on Trusting Trust", where he modified a compiler to insert backdoors in the compiled binary even when the source code was not containing anything malicious. This attack established that the entirety of a software system depends on the unbroken chain of trust.

Following Thompson's work, supply chains remained mostly theoretical or small operations, however the landscape changed with the rise of open-source package distributors and their ecosystems. A new attack surface of extreme scale was introduced, modern software extensively relies on third-party libraries and dependencies, dependencies having their own dependencies which creates a complex tree that can be exploited. Popular tools and software include hundreds if not thousands of transitive dependencies, each one of them can be a potential attack vector for malicious code injection, new attacks consisting of typosquatting, malicious build files, code injection and more.

2.1.2 Landmark Supply Chain Incidents

Multiple high profile supply chain attacks have been deployed and demonstrated the high risk of this threat vector. Table 2.1 provides an overview of these incidents.

The SolarWinds attack in 2020 is a state sponsored attack, which compromised the build system of SolarWinds network, by inserting a complex backdoor called SUNBURST into software updates, the attackers managed to deploy the malware to approximately 18.000 customers, including US government agencies and top grossing companies. The attack managed to stay undetected for over eight months and demonstrated how a single compromised dependency in the chain can provide access to a very large number of dependents.

Besides the SolarWinds attack, there was also the Log4Shell vulnerability, even though this was not a malicious attack, this critical vulnerability present in a extremely popular Java library affected millions of applications and threw the industry

Table 2.1: Summary of Landmark Supply Chain Incidents

Incident	Year	Type	Impact
SolarWinds	2020	State-sponsored	~18,000 customers including US government
Log4Shell	2021	Vulnerability	Millions of Java applications
event-stream	2018	Account compromise	Cryptocurrency theft via Copay wallet
ua-parser-js	2021	Account compromise	8M weekly downloads; crypto miners
xz-utils	2024	Long-term infiltration	SSH backdoor targeting Linux systems

into chaos for multiple days, requiring maintainers and organization to move quickly and patch their software.

The event-stream incident from 2018 is another example of a supply chain attack, where an attacker gained access to a maintainer account for the package 'event-stream', which is a popular npm package with over 2 million weekly downloads. Social engineering was used together with code injection targeting a cryptocurrency wallet called 'Copay Bitcoin wallet', with the goal being to steal cryptocurrency. The total amount of stolen money is unknown, and even though it is thought to be a small amount, this attack demonstrated yet again how supply chain attacks can be deployed.

In 2021, 'ua-parser-js', which is a nother popular npm package with approximately 8 million weekly downloads, was compromised in a similar way to the event-stream attack, by gaining access to a maintainer account the adversaries injected cryptocurrency miners and stealers, targeting users and organizations including Microsoft, Amazon and Meta.

A more recent attack, from 2014, is the 'xz-utils' backdoor, where an adversary spent multiple years building trust by contributing to the library, before inserting a complex backdoor in the test files and build system targeting SSH on Linux system, this attack was complex and it contained multiple accounts mentally draining the maintainer. The backdoor was discovered by chance, luckily before being merged into major Linux distros, by Andres Freund, a software engineer working at Microsoft, by observing that on a Debian system ssh logins were unusually slow compared to previous versions, OpenSSH is interacting with a part of xz-utils.

2.1.3 Rust-Specific Supply Chain Incidents

The Rust ecosystem is still developing compared to other ecosystems like npm and PyPI, however it still suffered from multiple supply chain attacks. Table 2.2 summarizes these incidents.

In 2022, researchers from SentinelOne discovered a supply chain attack which they named CrateDepression, the attack targeted GitLab CI pipelines through a typosquatting attack, by publishing a crate named 'rustdecimal', which is meant to be confused with 'rust_decimal', which is a legitimate crate providing decimal number support without rounding errors compared to Rust's native floating point

Table 2.2: Rust-Specific Supply Chain Incidents

Incident	Year	Attack Type	Target
CrateDepression	2022	Typosquatting	GitLab CI pipelines
faster_log / async_println	2025	Combosquatting	Cryptocurrency wallet keys
CI/CD targeting crates	Various	Env. exfiltration	GitHub/GitLab CI secrets

types (f32 and f64). The attack detected if the crate was used in a GitLab CI pipeline and would download and execute a binary written in Go, known as a language commonly used for malware since some antimalware solutions struggle to analyze newer languages, and Go is mature enough to be used for complex malware. This attack exfiltrated secrets from the environment and was able to persist in the system by modifying a configuration YAML file. Attacks like this one demonstrate how Rust’s memory safety and other security features are unable to protect against these kinds of attacks.

A more recent attack, from 2025, consists of multiple malicious crates including ‘faster_log’ and ‘async_println’, which was one of the motivations of this paper. These crates that targeted cryptocurrency systems developers contained code which exfiltrated Solana and Ethereum wallet keys from popular wallets present in the system by scanning for files and environment variables and sending the found data to C2 servers. This is not an example of typosquatting, since there are no similar legit crates, however they promised improved speed compared to the log crate and the println macro from the Rust standard library. These attacks were probably meant to be large-scale due to Rust being popular in cryptocurrency development, with projects including Solana being written in Rust.

Other attacks which are oriented on CI/CD exfiltration, where multiple Rust crates were identified to scan for environment variables related to CI/CD including ‘CI’, ‘GITHUB_ACTION’, ‘GITLAB_ID’ and others. These crates would exfiltrate environment files, secrets from the workflow and other credentials.

Multiple campaigns have been identified by security researchers, having dormant crates which contained obfuscated malicious payloads which would be executed at a later time when specific conditions were met.

2.1.4 Classification of Supply Chain Attacks

Several efforts were made by researchers to categorize supply chain attacks to have a deeper understanding of vectors and defensive techniques, which include:

Ohm et al. (2020) “Backstabber’s Knife Collection” This is a complex review of supply chain attacks which introduced a taxonomy to have them classified:

Typosquatting consists of publishing packages with similar names to popular packages, to deceive developers into installing them instead of the proper ones.

Dependency confusion exploits some package manager’s resolution to prioritize public packages instead of internally used ones. These packages have the same names as internal packages, where the names could be exfiltrated by

leaked source code, leaked documentation, guessing common names or internal information from employees.

Compromised accounts where by taking over legitimate maintainers and developers accounts through social engineering, credential leakage, phishing or other techniques, adversaries disguise themselves as a trusted contributor and inject malicious code, possibly months down the line after gaining access.

Malicious contribution similar to compromised accounts, but instead of gaining access or taking over an existing account, they submit malicious contributions to the project, sometimes being hidden behind a legitimate possibly large pull request, making use of complexity and human error.

Ownership transfer in some ecosystems it is possible to gain access to a package by being added as a maintainer or requesting ownership transfer for abandoned packages.

Ladisa et al. (2023) “A Taxonomy of Attacks on Open-Source Software Supply Chains” was published at IEEE S&P 2023. This paper proposed another taxonomy based around the lifecycle of software development:

Development phase attacks targeting local environments, such as IDE plugins, development tools and more.

Build phase attacks which compromise build systems such as CI/CD or build dependencies.

Distribution phase attacks which target package registries, distribution channels and infrastructure and more.

The taxonomy identified over 100 attack vectors across these phases and had defenses mapped to each one of the vectors, which revealed gaps in coverage, especially present in the build phase.

2.1.5 The Dependency Problem at Scale

We can consider the fundamental challenge of supply chain attacks to be the transitive nature of dependencies. Zimmermann et al. (2019) “Small World with High Risks: A Study of Security Threats in the npm Ecosystem” analyzed the npm ecosystem and revealed the following:

- The average npm package trusts 79 packages.
- Up to 40% of packages depend on at least one package with a public known vulnerability, even if it’s exploitable or not in a real world scenario.
- A single compromised maintainer account could affect multiple million packages through transitive dependencies.

Comparable comprehensive studies have not been done for crates, however the ecosystem seems to have a similar pattern. Rust's culture highly encourages code reuse and the development of small focused crates, similar to the Unix philosophy, which results in complex and deep dependency trees. A typical rust application can have hundreds of transitive dependencies.

Many crates suffer by the having a single maintainer, which is almost always an unpaid volunteer. The previously mentioned sz-utils incident demonstrated how attackers are able to mentally drain maintainers with harassment and spend multiple years in gaining trust before executing an attack.

2.2 Package Ecosystem Security Analysis

2.2.1 Vulnerability Databases and Advisory Systems

The detection and communication of known vulnerabilities is managed through databases maintained by diverse ecosystem-specific organisations, however maintainers are not always properly notified and incentivised to update dependencies.

The RustSec Advisory Database is managed by the Rust Secure Code Working group. It is the source of security advisories for the crates ecosystem, they cover vulnerabilities, including memory safety issues by poor unsafe code, cryptographic problems and logic errors. They include:

- Bugs that violate Rust's safety guarantees, even when not directly exploitable.
- Unmaintained crates which usually accumulate security issues.

Each advisory includes metadata, such as affected versions, patched versions, a scoring system and links to relevant CVEs.

The Common Vulnerability Scoring System (known as CVSS), provides a standardized method for rating vulnerabilities based on severity and impact, the score is from 0.0 to 10.0, being calculated from eight metrics:

Attack vector describes how the vulnerability can be exploited, such as network, adjacent, local and physical.

Attack complexity describes the required conditions for the attack.

Privileges Required describes the level of access the attacker needs.

User Interaction if the attack requires any user interaction such as clicking a button or opening a file.

Scope describes if the problem affects anything outside of the component.

Confidentiality Impact measures the impact on data confidentiality.

Integrity Impact such as data tampering, code execution and others.

Availability Impact such as Denial of Service and others.

The mathematical formula for the score combined these metrics into the score, which is furtherly classified as none, low, medium, high and critical. The formula is as follows:

Cargo audit is the official tool for checking a project's dependencies against the RustSec database. Using a command line interface, it includes features such as:

- JSON output for automation and further integration.
- Warnings for unmaintained crates.

Even though cargo audit is a great tool and highly recommended, the adoption in CI/CD is lacking. A 2023 survey found that fewer than 30% projects using CI/CD included cargo audit checks.

2.2.2 Large-Scale Ecosystem Security Studies

Multiple research efforts were already conducted on some package ecosystems, however the Rust ecosystem is still underexplored as it is still in its early stages compared to ecosystems like npm and PyPI, which have been studied.

The npm ecosystem has been studied in multiple papers:

- Staicu and Pradel (2018) analyzed over 235.000 npm packages for injection vulnerabilities, finding thousands of exploitable packages.
- Zahan et al. (2022) examined the relationship between maintainer activity, contributor diversity and security findings.
- Lauinger et al. (2017) analyzed client-side libraries in websites and found extremely high usage of outdated and vulnerable versions.

The PyPI has also received attention from researchers:

- Vu et al. (2020) documented typosquatting attacks in PyPI, which proposed detections.
- Guo et al. (2023) analyzed malicious package campaigns and investigated evasion techniques.
- PyPI infrastructure has implemented complex malware detection including YARA rules and behavior analysis.

Cross-ecosystem studies have been also been done:

- Decan et al. (2018) compared npm, CRAN and RubyGems
- Ponta et al. (2020) compared patterns between Maven, npm and PyPI.

However, despite Rust's high growth and importance in the industry, including security critical fields, there is a lack of a comprehensive security analysis of crates.io, existing research is mainly focused on memory safety and other language features, but not on the entire ecosystem. This thesis aims to address this gap.

2.3 Typosquatting in Package Registries

2.3.1 Definition and Threat Model

Typosquatting is a form of attack where adversaries register malicious packages with names similar to popular packages, usually containing just a typo, this is to exploit human error, in case developers mistype packages when adding dependencies.

A possible threat model could be:

Attacker capabilities Register arbitrary package names, publish malicious code and track downloads.

Attack surface All packages installation events (manual editing of Cargo.toml, cargo add or automated tools).

Victim behavior Occasional human errors, poor verification of dependencies, reliance on ai autocomplete.

Attacker goals Code execution, which can lead to credential theft, cryptocurrency theft and mining, further remove code execution and more.

2.3.2 Typosquatting Patterns

Researchers have identified several patterns used in typosquatting attacks, such as:

Character substitution Replacing characters with similar alternatives, such as 0 instead of o, l instead of 1, i instead of l and others. This is not as common due to user's not usually copy-pasting package names, and being less effective on getting users to install the dependencies.

Character omission By removing a single character, such as 'rust_decimal' being typosquatted by 'rustdecimal', this being a previously mentioned attack.

Character addition Similar to character omission, but instead of removing a character another one is added.

Character transposition by swapping adjacent characters, such as 'tokoi' instead of 'tokio', this is another easily spotted pattern, however might get some users due to users typing behavior.

Separator manipulation by changing, adding or removing separators, such as 'serde-json' vs 'serde_json', this is a common pattern in the rust ecosystem due to the common use of underscores and dashes.

Prefix and suffix manipulation by removing or adding prefixes and suffixes, usually common ones for the package ecosystem, such as 'rust-serde' vs 'serde'.

Combosquatting which is combining a legitimate name with commonly used modifiers, such as 'faster_log' and 'async_println' which are previously mentioned.

2.3.3 String Similarity Algorithms for Detection

Typosquatting detection mostly relies on calculating the similarity between multiple package names using different algorithms, such as:

Levenshtein Distance also known as edit distance, which counts the minimum single character operations to modify a script to another.

Damerau-Levenshtein Distance which is an extension of the Levenshtein distance, which adds transposition of adjacent characters, this is useful because of the keyboard layout permitting these kinds of errors.

Jaro-Winkler similarity it computes the number of matching characters, number of transpositions and computes a score combining these factors, after this it adds extra weight to some common prefixes.

Keyboard proximity distance which measures the physical distance between keys on a keyboard, usually QWERTY because of its dominance.

Phonetic similarity there are algorithms like soundex or metaphone which compare how names sound rather than their spelling, however this is not very useful for typosquatting detection, since usually package names are transmitted through text, not audio.

There is also the possibility of running all of these algorithms and computing a meta-score based on weights on the previously mentioned algorithms or others.

2.3.4 Typosquatting in the Rust Ecosystem

Compared to other ecosystems, Rust's crates.io has a more specific naming pattern that affects typosquatting attacks and detection. There is a heavy use of hyphens and underscores, together with common suffixes like '-rs' and others.

Rust Foundation Typosquatting Protection The Rust Foundation has implemented typosquatting detection using the crate 'typomania', it computes similarity scores for new crates against existing popular ones, and flags names very similar to current popular crates, requiring manual approval. It is confirmed that this system has successfully prevented multiple typosquatting attacks.

2.4 Malicious Code Detection in Package Ecosystems

2.4.1 Categories of Malicious Behavior

Malicious code from package ecosystem exhibits pattern such as the following:

Data exfiltration, which involves stealing information from the victim system, such as environment variables (API keys, credentials, database connection strings, sensitive settings), files (configuration files, SSH keys, cryptocurrency wallet keys, sensitive configuration files), browser data, clipboard contents and more. The crates

previously mentioned that targeted Solana and Ethereum users fit into this category, searching for cryptocurrency wallets and private keys.

Reverse shells and backdoors, establishing persistent access to the victim system.

Cryptocurrency mining, mostly CPU-based mining of Monero due to its privacy and hard to scale mining algorithm.

Payload downloading, by downloading additional malware. It sometimes uses legitimate services as the hosts such as github, free CDNs, pastebin or others.

Conditional execution, sometimes, the malware authors implement malicious behavior execution only in some branches of the code, ran only if specific conditions are met, such as a CI environment being detected, executing only in specific countries or after a specified date.

2.4.2 Static Analysis Approaches

Static analysis is analysis of a binary or code without executing it, it is very scalable compared to hybrid analysis but more limited in its capabilities.

Signature-based detection relies on known malicious patterns such as having hardcoded specific IP addresses and domain names, characteristic code from known malware or specific obfuscation patterns used in previous malware.

YARA rules are widely used for signature-based detection, they allow us to define specific, even complex patterns combining features such as string matches, byte patterns and metadata.

Behavioral pattern matching looks at more complex patterns, such as sequences of calls, network connections from unexpected modules, file system operations on unusual files and more and more.

High entropy in code can also indicate either encrypted or compressed code, encoded or encrypted strings, large embedded binary data and more, usually not used in benign packages.

Dependency analysis is used to examine which dependencies a package has, we can identify if there are unexpected network calls or dependencies in a package that should not include them, having other dependencies on packages with known security issues or unusual version requirements, such as requiring an extremely old version of a package.

Metadata analysis, looking at new maintainers, recent ownership transfers, unusual releases such as multiple versions in a rapid succession, rewrites of old commits, changing version tag pointers to other commits, and others.

2.4.3 Dynamic Analysis Approaches

Dynamic analysis is the counter part to static analysis, that involves emulating the code or executing it in a sandboxed or controlled environment to observe the runtime behavior.

Sandboxed execution is running packages in isolated environments such as virtual machines or containers, with the interception of system calls, network connections and more to assess the behavior of the package.

Behavioral monitoring is used to record network connections, file accesses, process execution, environment variable access and more.

Differential analysis refers to comparing behavior between two versions, to check if there is any change in behavior such as new network connections to existing or new servers and other different patterns.

2.4.4 Machine Learning Approaches

Feature based classification consists of extracting features such as metadata from packages and training machine learning classifiers. It can be effective for some patterns. Code metrics, dependency graph, metadata and behaviors are analyzed.

Text based classification is run on the source code and applies NLP techniques such as bag of words, TF-IDF, code embeddings, sequence models and more to classify code.

Graph neural networks model dependencies as graphs and are also used for malicious code detection by analyzing for suspicious patterns.

Machine learning approaches can be effective, however they face challenges such as class imbalance, concept drift, adversarial evasion especially for text based classification and LLMs.

2.5 Build Script Security Risks

2.5.1 Build Script Security in Rust

Rust's `build.rs` feature is very powerful and allows crates to execute arbitrary Rust code at the compile time of a package and its dependents, cargo first checks for a file named `'build.rs'` in the root directory of a crate, compiles its code and executes it before compiling the main crate. This allows the build scripts to perform advanced and complex operations and possibly infect systems.

There are multiple legitimate uses of build scripts such as native library linking to C libraries, code generation such as FFI bindings, protocol buffer compilation and embedding resources, compile-time configuration such as checking for optional dependencies, setting flags based on the environment and current platform, calling external tools and more.

2.5.2 Security Implications

However, the same power that enables build scripts to be a great tool for the previously mentioned use cases, also makes it a very dangerous attack vector, enabling multiple security risks such as:

Arbitrary code execution; the build script has access to the filesystem, network, environment variables, process execution and more.

Attack surface expansion, due to the dependency trees that packages have, the attack surface doesn't consist of just the compiled package or its direct dependencies but also the entire dependency tree.

CI/CD environment targeting: Build scripts are powerful enough to be able to properly detect CI/CD environments using advanced detection that can consist of detecting environment variables, system privileges and information such as the CPU, OS, storage and installed software, existing or missing files and others.

Malicious build.rs code can be obfuscated and try advanced detection evasion such as conditional execution, and executing other payloads at runtime if specific conditions are met.

2.5.3 Community Discussions and Proposed Mitigations

The Rust community has discussed build.rs security multiple times, especially after the previously supply chain attacks.

Sandboxing was one of the proposed solutions, by restricting network and filesystem access, requiring explicit permissions and implementing user prompts.

However as of the time of writing this thesis in 2026, there is no sandboxing present, in part due to the complexity of such a system and the backwards compatibility concerns.

A recommendation is to minimize build.rs pollution, by avoiding using it when it's not necessary, and proper review of changes to build.rs scripts.

2.5.4 Heuristic Detection of Suspicious Build Scripts

Given the lack of sandboxing or other mitigations, heuristic detection can be considered a practical approach to identify malicious build scripts, we can look for patterns such as:

Network library detection, since most build script's shouldn't have the need of making network requests, the presence of HTTP libraries is unusual and be considered a "red flag" in cases where the build script shouldn't be making network connections.

Process spawning is commonly used for invoking tools, however some patterns can be suspicious such as spawning shells such as cmd, powershell, bash and others, executing downloaded binaries or decoding and executing obfuscated code from remote or embedded sources.

Hardcoded raw IP addresses in build scripts is unusual, and legitimate uses usually involve trusted domains and hostnames. The usage of raw IP addresses can suggest the usage of C2 servers or downloading malicious payloads.

Suspicious domain names, such as free or commonly abused TLDs are also suspicious, some TLDs have less strict registration requirements and enforcement.

High entropy strings are very suspicious and can indicate obfuscated and encrypted code or payloads.

Link directive analysis can be also used, since linking to libraries downloaded by the crate from a suspicious source can indicate malicious behavior.

2.6 Large Language Models for Code Analysis

2.6.1 Evolution of Code-Aware Language Models

The application of LLMs to code understanding as rapidly evolved in the past years, starting from models such as CodeBert, Starcoder To Qwen2.5-Coder and Claude Opus.

2.6.2 LLMs for Vulnerability Detection

Several researchers have evaluated LLMs on security tasks such as vulnerability detection.

The Li et al. (2023) "Assessing the Promise and Pitfalls of ChatGPT for Automated Vulnerability Detection" study has evaluated ChatGPT's older 3.5 and 4 models on vulnerability detection, finding that they have a reasonable accuracy on well known patterns, but that they struggle with uncommon vulnerability types and more complex code, and they have a high false positive rate.

"How Well Do Large Language Models Serve as End-to-End Secure Code Producers?" by Zhang et al. (2024) is a evaluation of multiple LLMs on SecurityEval, which found that SOTA LLMs still introduce vulnerabilities in generated code, and that they detect vulnerabilities far better than avoiding generating them at all.

2.6.3 Challenges in Large-Scale LLM Security Analysis

Using LLMs at a large scale such as ecosystem-wide level introduces multiple challenges.

Firstly, the context window is limited even for large context models (such as the 128K+ tokens models), which cannot process larger codebases in a single prompt. We can chunk code into batches, however this can lead to loss of context and has a significant accuracy impact and increased time and cost. We can also prioritize analysis of specific files such as build.rs or files imported in build.rs, however this is also not a perfect solution since malicious code can be hidden in other files or imported dynamically. RAG can also be used however it still has the problem of losing context, increased time and cost and less accuracy than fitting the entire codebase in the context window.

We talked about the computational cost a bit, however LLM inference is computationally expensive in all cases, even with batching and local hosting there still is a significant cost and time cost.

LLMs also have the tendency of hallucinating, which can lead to reporting false positive or false negatives, even changing the output format and conclusion based on small changes or different seeds, even with the temperature set to 0.

Prompt engineering can help with accuracy however it isn't a complete and perfect solution, and there is no real way of finding the best prompt, longer prompts also occupy more of the context window.

2.6.4 Local vs. API-Based LLM Deployment

For an ecosystem scale analysis, deployment architecture is key.

API-Based deployment is easier to set up and has no maintenance cost, however it has a cost due to the sheer scale.

Local deployment has the advantages of being cheaper for the inference, however it has a maintenance and investment cost for the hardware.

For this research, we are using a small local deployment of Qwen2.5-Coder-14B, providing privacy and 0 cost per query besides the already owned hardware and electricity cost.

2.7 Secret Detection in Source Code

2.7.1 The Problem of Hardcoded Secrets

Hardcoded secrets represent an ongoing security problem especially in novel packages and software. Secrets such as API keys for cloud providers and services, database connections, private ssh keys, tokens, jwt secrets and others can be accidentally included in commits to the source code.

Developers move fast, sometimes hardcoding during development and forgetting to remove files and committing them to version control.

2.7.2 Secret Detection Approaches

Regular expressions can be treated as a form of signature based detection, using patterns for known secret formats such as Github tokens, private keys, and cloud provider API keys can be an effective way of detecting secrets.

Like encrypted and obfuscated strings, secrets usually have high entropy compared to normal code.

Machine learning approaches can also be an useful tool for secret detection, by training models on known past codebases with revoked secrets can be effective.

2.7.3 Gitleaks

Gitleaks is the industry standard for open source secret detection in git repositories, it uses a rule based approach such as regular expression patterns, entropy thresholds, path patterns, keywords and others. There are over 150 rules present. Gitleaks provides full repository scan, being able to scan all the commits and it can also provide CI/CD integration. It is also a great tool to implement due to it's JSON output, it is easy to implement into existing frameworks, pipelines and infrastructure. Written in go, it is highly performant even on large repositories.

2.7.4 Secrets in the Rust Ecosystem

Rust's crates have permanent publication, they cannot be completely removed, only tyanked, which remain downloadable if they version is set before the point of yanking. Developers may accidentally include configurations files, environment files, debug code with secrets and more.

2.8 Existing Tools and Related Work

2.8.1 Rust-Specific Security Tools

Cargo audit is the official RustSec vulnerability scanning tool it requires just the Project's Cargo files and lists all the known vulnerabilities of all of the dependencies. It is well maintained and accurate.

Cargo deny is a plugin that enforces dependency policies, such as license checking, advisory checking (similar to cargo audit), bans on specific crates or versions, duplicate detection and source verification.

Cargo crev is a code review system based on web of trust, maintainers publish signed reviews of crate versions and users are able to configure manually trust settings for reviewers, however it has a low adaptation, relying on a small community of volunteers to manually review packages.

Cargo vet is Mozilla’s supply chain security tool, where organizations publish audits to prove a crate’s safety. It focuses on organizational trust compared to our need of automated analysis.

Cargo tree let’s us visualize the dependency tree of a project, showing direct and transitive dependencies, however it lacks security analysis features.

2.8.2 General-Purpose Security Tools

Gitleaks as described in section 2.7.3, is the standard for open source secret detection in git repositories.

Semgrep is a fast open source static analysis tool, supporting manu languages including Rust, having a very large rule library, support of custom rules and pattern matching with semantic awareness.

Dependabot has automated dependency updates and can create pull requests for vulnerable dependencies, and can be integrated with Github and GitLab, however it is not a malware detection tool.

2.8.3 Commercial Supply Chain Security Platforms

Socket.dev provides supply chain attack detection of npm and PyPi packages, meta-data anomaly detection and network call analysis, however it lacks Rust support as of now.

Phylum provides supply chain risk assessment with experimental Rust support, machine larning risk scoring, author reputation analysis and dependency relationship analysis.

2.8.4 Academic Research on Ecosystem Security

”Understanding Security Threats in Open-Source Software Packages” (Survey papers) provides surveys of attack techniques and defenses, analysis of detection effectiveness and categorization of attacks

”On the Prevalence of Vulnerable Dependencies in the npm Ecosystem” analyzed large scale data on npm.

”A Measurement Study of Security Threats Collected from npm and PyPI” analyzed both npm and PyPI ecosystems, including typosquatting and dependency confusion.

2.9 Distributed Systems for Large-Scale Analysis

2.9.1 Challenges of Ecosystem-Scale Analysis

Analyzing an entire package ecosystem has multiple, highly-complex challenges including:

The scale of the ecosystem, crates contains over 150.000 crates with each having multiple versions, easily spanning over a million releases, analysis would involve downloading each repository, running multiple advanced modules, storing results and managing temporary storage for execution speed.

Dependency ordering, transitive risk propagation is a key feature and it requires analyzing dependencies before the dependents. This heavily complicates parallelization.

Analysis time is very variable, some crates taking seconds to analyze others can take minutes for large codebases, spanning into hours when scanning the entire commit history with LLMs.

Fault tolerance is also a challenge, the ecosystem has multiple broken repositories, failures such as network issues, timeouts and others are common and we must recover gracefully without losing progress.

Resource management is also required, concurrent analysis require disk space management due to the usage of a ramdisk, memory usage is also high together with network bandwidth, cpu usage and database connections.

2.9.2 Work Queue Architectures

A distributed work queue architecture is a common solution for analysis similar to this, we have multiple options for queues.

Redis provides primitives for implementing work queues, 'BRPOPLPUSH' is a command for atomic claim and move operations, it also provides features such as hashes for storing metadata.

RabbitMQ is a message broker providing persistent work queues, message acknowledgement, retry possibilities, load balancing and more.

Apache Kafka provides high throughput, partitioning and more advanced features, however it is more complex to set up and can be overkill for this use case.

For this research, we picked Redis due to familiarity, simplicity performance and native atomic operations which let us implement a simple work queue.

2.10 Summary and Research Gap

The reviewed papers show that even though supply chain security has started to receive more attention from researchers, there are still significant gaps:

2.10.1 Ecosystem-Specific Gaps for Rust

The rust ecosystem has received far less attention from researchers compared to ecosystems like npm or PyPI despite of it's growth in importance, especially in critical infrastructure and fields. No existing work combines vulnerability scanning, build script analysis with secret detection and LLM analysis into a unified comprehensive intuitive solution. Especially, due to build.rs scripts being an extremely powerful feature it is needed to have more research on how to protect upstream dependents.

2.10.2 Methodological Gaps

LLMs have been evaluated for vulnerability detection, however there is a lack of research on deployment at scale on an ecosystem level. Existing tools such as cargo audit work on individual projects, not on the entire ecosystem to provide transitive risk propagation and other more advanced features. Security tools are used in a reactive way, isolated, there is a benefit to correlate findings across multiple "dimensions" (vulnerabilities, secrets, build.rs behavior). Which might be useful for prioritization and revealing more complex patterns.

2.10.3 Contributions of This Thesis

The thesis addresses these gaps by designing a complex, comprehensive optimized framework for analysis integrating multiple modules with a focus on the Rust crates ecosystem, by implementing a distributed architecture, with an in memory work queue, web dashboard and a database that can be used for further analysis. The framework uses multiple modules such as heuristics build scripts analysis together with existing tools, we are conducting a large scale empirical assesment to more easily identify possible malicious packages and signal possible vulnerabilities. The plan is to open source all of the code, with the possibility of having the Rust foundation integrating parts of it or a modified version into their infrastructure, to provide better security for the entire ecosystem.

Chapter 3

Methodology

3.1 Research Approach

This thesis implements a novel framework with empirical evaluation of its application to the Rust crate ecosystem. The research process has 4 steps.

1. Problem identification and requirement analysis: literature review of supply chain security and attacks, identification of research gap and objectives establishment.
2. Design and development: Iterative implementation and design of the preset-ned framework, keeping in mind the requirement of scalability.
3. Empirical application: Execution against a subset of the crates.io database.
4. Evaluation and discussion: Result analysis, measurements, comparison and future work ideas.

3.2 Threat Model

3.2.1 System Model

The ecosystem has the following principal entities:

- Crate maintainers: Individuals/organizations that develop the crates. They control the repository, source code and metadata.
- Consumers: Developers who have at least a crate dependent on a project.
- Crates.io registry: The centralized repository that stores the crate versions, metadata and sometimes code, and provides them to the consumers.
- Build infrastructure: Cargo, rustc and CI/CD systems.
- End users: Users of software build using crates from crates.io, who can also be affected by any compromise in the supply chain

3.2.2 Attacker Model

We consider attackers with the following skills: Package registration - an attacker can register any crate name allowed by crates.io and with no existing crate Code publication - an attacker can publish crate versions containing arbitrary code Build-time execution - Code can be executed at compile-time using build.rs Metadata manipulation - Keywords, descriptions and documentations from a crate page are written by the authors Ecosystem knowledge - knowledge of popular crate names,

popular keywords, types of crates, naming conventions, etc. Resources - limited resources compared to nation sponsored actors

We assume attackers cannot modify crates infrastructure or compromise its server, or find any exploits that would allow them to modify crates owned by other maintainers without having direct access to their accounts, or do Man in the Middle (MITM) attacks between victims and crates.io.

3.2.3 Attacker Classes

Class	Motivation	Resources	Examples
"Skid"	Experimentation	Low	Copy-paste/Vibecoded malware
Criminal	Financial gain	Medium	Professionally developed Stealers, C2, ransomware
Advanced Criminal	Targeted financial gain	High	Sophisticated attacks, targeted espionage
Nation-State	Espionage, sabotage	Very High	Long-term infiltration (i.e. xz-utils), infrastructure compromise

The framework primarily targets "skids" and "criminal" attackers, whose techniques are less advanced and more easily detectable using static analysis. Advanced attackers could evade automated detection due to factors like complex obfuscation and understatement of the system.

3.2.4 Attack Vectors Addressed

Attack Vector	Detection Module	Detection Approach
Known vulnerabilities	Cargo Audit	RustSec database matching
Hardcoded secrets	Gitleaks	Regex + entropy analysis
Malicious build.rs	Build Script Heuristics	Network, IP, TLD, process, entropy patterns
Typosquatting	Typosquat Scorer	Multi-algorithm string similarity
Obfuscated/novel malware	LLM Analysis	Semantic code understanding
Transitive compromise	Dependency Propagation	Graph-based risk propagation

3.2.5 Scope and Assumptions

We operate the framework under these assumptions:

1. The source code is available and hosted on a git repository that doesn't rate limit or restrict our access, grabbed from the crates.io database.
2. LLM reliability, due to models becoming better at code reviews, and the framework being easy to extend for API usage with state of the art models
3. Attackers reuse patterns and techniques that can be found through the use of heuristics.

The following are out of scope for this thesis:

- Dynamic/Runtime analysis using sandboxed or emulated execution
- Analysis of compiled binaries
- Analysis of network behavior at runtime (included in runtime analysis)
- Compromised crates through hacked accounts or infrastructure exploits

- Detection of social engineering

3.3 Data Collection Strategy

The crate metadata source is the crates.io database dump published by the crates.io team, containing CSV files of all published crates including metadata, such as name, description, repository url, etc; version information, download statistics, and dependency relationships

These CSV files are processed and inserted into a PostgreSQL database using a custom pipeline that resolves dependency references and constructs a dependency graph. The source code of the repositories is obtained by cloning the last commit of the git repository from the crate metadata.

The system supports a configurable scope using the 'use_only_first_x_crates' field from the configuration toml file, together with the single crate analysis parameter, allowing the analysis of a single crate and its dependencies.

3.4 Analysis Framework Design

The design was focused on the following principles:

- Modularity: Having independent modules that can removed or added easily, possibly requiring DB changes
- Horizontal scalability: We have the possibility of adding worker nodes from a single or multiple machines to increase throughput
- Topological processing: Processing dependencies before dependents to enable transitive risk propagation
- Fault tolerance

3.5 Evaluation Metrics

The evaluation will consist of two parts:

Security analysis

- Number and severity distribution of vulnerabilities
- Number, types, average/median entropy of secrets
- Characteristics of build scripts
- Distribution of LLM scores
- Number and quality of typosquatting candidates
- Transitive risk propagation statistics

System performance

- System Throughput
- Worker speed-up scale
- Resource utilization

Chapter 4

System Architecture

4.1 High-Level Architecture Overview

The framework uses a producer-consumer architecture with a central layer, the main components are:

1. Data Pipeline - Imports all the crate information required from the crates.io database into PostgreSQL
2. Producer - Constructs the Redis dependency graph and inserts crates with no unprocessed dependencies into the queue
3. Worker - It claims a crate from the Redis ready queue, runs all modules and inserts the results into PostgreSQL
4. Final processing: Computes transitive dependency risk after all analyses are finished.
5. Web Panel - A browser interface for visualizing results




Figure 4.1: High-level architecture of the distributed crate analysis framework.

4.2 Data Ingestion Pipeline

This pipeline processes CSV files from the crates.io database dump archive:

- crates.csv - contains core metadata (ID, name, repository URL, etc.) - crate_downloads.csv - download count for each crate - versions.csv - version specific metadata - dependencies.csv - crate dependency associations between versions, including the dependency type and version id

These CSVs are all processed, firstly we resolve crate-version relationships and after processing we perform batch inserts into PostgreSQL.

4.3 Dependency Graph Construction

The dependency graph is a DAG where nodes represent crates and edges represent dependency relationships. The graph is build by:

1. Query PostgreSQL for all dependencies
2. For each crate build a dependency set in Redis, tracking unprocessed dependencies.
3. Build a reverse index to be able to get which crates depend on a target crate.
4. Pushing crates ready to process (with zero unprocessed dependencies) into the Redis ready queue.

After a worker analyses a crate, it has to atomically update the dependency sets of all that crate's dependent crates. If any of these's sets becomes an empty set, it is pushed into the ready queue.

Dependency cycles are detected using depth first search and are resolved by breaking cycle edges.

4.4 Distributed Work Queue

The Redis queue has the following structures:

Redis Key	Type	Purpose
crate:ready	List	FIFO queue of crate IDs ready for analysis
crate:in_progress	Set	Crate IDs currently being analyzed (atomic claim)
crate:completed	Set	Crate IDs successfully analyzed
crate:failed	Set	Crate IDs that failed analysis

A worker uses 'BRPOPLPUSH' to do an atomic claim of a crate, this ensure there are no race conditions. Failed crates are marked as failed but are not requeued and do not block dependent crates from being processed. This can be further improved if

4.5 Database Schema Design

The PostgreSQL database contains the following tables:

Table	Description	Key Fields
<code>crates</code>	Crate metadata from crates.io	id, name, repository, crate_downloads, timestamps
<code>dependencies</code>	Dependency edges between crates	crate_id, dependency_id (composite PK)
<code>scan_results</code>	Aggregated analysis results per crate	FK to crates, LLM score, boolean flags, build.rs heuristics, entropy
<code>cargo_audit_results</code>	Individual vulnerability findings	UUID PK, FK to crate, rustsec_id, CVSS score
<code>gitleaks_results</code>	Individual secret findings	UUID PK, FK to crate, rule_id, secret, location, entropy
<code>typosquat_results</code>	Potential typosquatting pairs	FK pair to crates, multiple similarity scores + combined score
<code>runner_metadata</code>	Analysis run tracking	timestamp PK, last checked crate

The schema was designed to use indexing for commonly used queries for filtering, combined scores and crate ids.

4.6 Web Dashboard

The web dashboard is implemented as a standalone project using Axum and provides an API for querying the database. The frontend supports the following features:

- Dashboard view - Showing high level statistics such as crates analyzed, vulnerability count, secrets found and others providing charts for easy visualization.
- Crates browser - a paginated, sortable and filterable list of all the analyzed crates.
- Scan results view - A filterable, sortable and paginated list of all crates with their metadata.
- Scan results page - Same as the others in terms of feature, showing information such as build.rs heuristics, vulnerability results and more.
- Vulnerability viewer - Lists all cargo audit findings for all crates with the serverity score and links to the advisory page.

- Secrets viewer - Shows all gitleaks findings for all crates, with rule IDs, entropy and location in the repository.
 - Typosquatting viewer - Shows all typosquatting pairs with all of the similarity scores from all algorithms used, together with the combined score
 - Dependency browser - Shows the dependencies between crates.
- All endpoints support pagination, sorting and multiple filtering.

4.7 Deployment and Infrastructure

The workers are containerized using Docker to ensure easy scalability on a single or multiple machines, the sample production deployment is the following:

- PostgreSQL database with persistent storage
- Redis for the work queue and dependency graph
- Worker containers, by using a custom Dockerfile we can scale multiple workers using 'ocker-compose up -d --scale worker=N' where N is the number of workers to start. Each worker has a limit of 2 cores and 4GB of RAM.

To achieve maximum I/O performance the analysis is done on a temporary mounted RAM disk at /tmp/crates, eliminating the disk I/O bottleneck during cloning and analysis.

Multiple machines can be used by configuring the workers on multiple machines to connect to the central Redis instance, this enables scaling across a cluster if needed.

Chapter 5

Analysis Modules

5.1 Cargo Audit - Vulnerability Scanning

The cargo audit module uses cargo audit, which is integrated with the RustSec advisory database, a maintained database of security advisories for crates. For each crate's local git clone of the latest hash the module executes the command 'cargo audit -json -no-fetch', the output a JSON from which we can extract vulnerability information and severity scores.

Severity scores are calculated by implementing the CVSS v3.1 calculation from the raw vector strings. Scores are normalized to a 0-100 scale for comparison.

Cargo audit cannot run locally in parallel due to a shared lock, to mitigate this we used docker containers for parallelization.

Results are stored in the 'cargo_audit_results' table and 'scan_results' table for aggregation.

5.2 Gitleaks - Secret Detection

The Gitleaks module detects hardcoded secrets, API keys, tokens, and other credentials present in the source code. Gitleaks output JSON is parsed to extract the following information:

- Rule ID - which represents the type of detected secret
- Secret - The detected value
- Location - Path and line number of the secret from the repository
- Entropy - Calculated shannon entropy by gitleaks of the string, which can indicate if it's a real secret or a false positive

The results are stored in the 'gitleaks_results'. The module is prone to returning false positives due to test secrets present in most of the cases.

5.3 Build Script Analysis

The build script analysis module has six heuristic implementations to detect potential malicious behavior:

- Network call detection - Strings related to network libraries and calls such as 'reqwest', 'hyper', 'curl', 'ureq', 'tokio::net', 'std::net::TcpStream' and others. Net-

work calls, as previously mentioned, can be used to download payloads or exfiltrate data.

- Link directive detection - Using Regex, we look for patterns such as link attributes, mostly used for FFI and non malicious behaviors, these can be used to link to malicious libraries.

- Shannon entropy calculation - High values can indicate obfuscated or encoded data and code, we also compute an average entropy of all '.rs' files from the repository to compare against.

- Process spawning detection - Using Regex, we look for patterns such as 'Command::new()', '.exec()', '.spawn()' and others. Commonly used to invoke external tools, however it can also be used to execute payloads, attackers often pair this with network calls to download the malicious payload.

- Raw IP addresses - Regex based detection of IP addresses, we exclude local and private ranges, however non private and non local IPs usually are an indicator of suspicious behavior instead of using domain names.

- Suspicious TLDs - Using Regex, we look for free and commonly used TLDs by attackers, such as '.tk', '.ga', '.cf' and others. These TLDs are often used due to them being free or having a low cost of rental.

5.4 LLM-Assisted Malicious Code Detection

The LLM analysis module uses a locally hosted LLM using LM Studio to perform analysis of crate source code. This module is optional due to extremely high computational cost compared to the rest of the modules. It operates in the following way:

- Collects all source code files from a target extension list
 - Source files are chunked to batches that fit into the model context window. Files are prioritized by relevance such as 'build.rs', 'lib.rs', 'main.rs' and 'mod.rs'
 - Each batch is sent to the LLM together with a prompt that instructs the LLM to evaluate the code for potential malicious behavior, such as reverse shells, unauthorized network access, data exfiltration, cryptocurrency mining and more.
 - The LLM returns a score from 0 to 100 together with a summary for each batch. The final score is set to the maximum score across the batches.
 - Crates with a score above the set threshold are flagged for human review.
- The module has a five minute timeout per batch.

5.5 Typosquatting Detection

The typosquatting detection module computes pairwise name similarity between crates in parallel using multiple string similarity algorithms:

Name	Weight	Notes
Levenshtein Distance	20%	Minimum number of single-character insertions, deletions, and substitutions required to transform one string into another.
Damerau-Levenshtein Distance	25%	Levenshtein with consideration of character transpositions
Jaro-Winkler Similarity	25%	Computes character-level matching with a bonus for common prefixes.
Keyboard Proximity Distance	15%	Based on the physical distance between characters on a QWERTY keyboard layout, using Euclidean distance between key positions.
Prefix Similarity	15%	Measures the length of the common prefix relative to the string lengths, with a bonus for prefixes of 3 or more characters.

The combined score, after applying the weights, is normalized to a 0-100 scale.

5.6 Transitive Dependency Risk Propagation

After the complete analysis, another step computes the 'has_malicious_dependencies' flag for each crate. The computation is done via a single SQL UPDATE query.

Chapter 6

Implementation Details

6.1 Technology Stack

Component	Technology	Justification
Core application	Rust	Memory safety, performance, async and concurrent support
CLI framework	Rust + Clap	Easy command definition with subcommands
Async runtime	Tokio	Industry-standard Rust async runtime, speed and ease of use
Database ORM	Diesel (async)	Type-safe SQL queries, fastest orm, migration management
Connection pooling	bb8	Fastest async connection pool for PostgreSQL
HTTP client	reqwest	Easy async HTTP for LLM API communication
JSON parsing	sonic_rs	High-performance JSON parser (faster than serde_json)
Work queue	Redis	Minimal latency, atomic operations, Lua scripting
Data storage	PostgreSQL	Fast, Relational integrity, complex query support, JSONB
Web framework	Axum	Async, Ease of use, performance, Tower middleware ecosystem
Parallelism	Rayon	Easy parallelism for typosquatting analysis
Logging	tracing + tracing-subscriber	Structured logging with file and console output, helpful for debugging and information
Containerization	Docker + Docker Compose	Reproducible and easy deployment, horizontal scaling

6.2 CLI Application Structure

The CLI app has the following CLI commands and structure:

Command	Description
update-database	Downloads and imports the crates.io CSV dump into PostgreSQL
run-analysis	Constructs dependency graph and populates Redis queue
worker	Starts a worker process that consumes from the queue, used from docker
compute-dependencies	Post-processing for dependency risk flags
queue-status	Displays Redis queue statistics
recover-stale	Recovers items from crashed workers
clear-queue	Resets the entire queue state
fix-cycles	Detects and resolves dependency graph cycles
scan-crate --name X	Single-crate analysis with full dependency tree
queue-crate --name X	Queues a single crate and its transitive dependencies
check-llm	Verifies LM Studio connectivity
typosquat	Runs typosquatting detection on the entire database

6.3 Redis Queue Implementation

The redis queue implementation was designed to handle multiple complex scenarios:

- Topological ordering of crates to ensure the dependents are analyzed first
- Atomic operations - Atomic claim of crates, handling crate completion, dependent updates and others happen in a atomic manner to prevent race conditions
- Blocking pop - BRPOP is used to avoid busy waiting
- Statistic tracking - The queue has counters for total crates, processed crates and others to provide monitoring capabilities.

6.4 Cycle Detection and Resolution

Dependency cycles can prevent crates from ever being processed as dependency cycles will cause crates to always have dependencies.

- Cycle detection - Using depth first search, we identify cycle edges
- Cycle breaking - Removal of edges allows crates to be processed, 'fix-cycles' is used for simple A-B-A cycles, and 'break-cycles' is used for more complex cycles.

6.5 Worker Lifecycle and Fault Tolerance

Each worker is an independent instance with a lifecycle that does;

- Initialization - Connects to PostgreSQL and Redis and sets up the state and handlers
- Main loop - Claim a crate from the ready queue in an atomic way, download its depository and run each module, after that insert the results into PostgreSQL and remove the crate from Redis.

- Shutdown - handle signals like SIGINT and SIGTERM by finishing the current crate and only after that exiting.

- Failed analysis produce placeholder results.

6.6 Batch Processing and Database Optimization

- Batch updating - INSERT and on conflict UPDATE statements for results
 - Chunked imports - The database population from the CSV files is done in chunks to improve performance and reduce memory usage
 - Connection pooling - bb8 is used to manage database connections
 - Indexing - Indexes on commonly used queries in the codebase are used, they do not increase the database size enough to be a problem

Chapter 7

Results and Evaluation

7.1 Dataset Overview

The analysis was done against a snapshot of the crates.io database obtained on [DATE] (atm 30 March 2026). NOTE: I will run the tool on a more recent version and will let it reach 100% or close to it. This is a "demo" for this class.

Metric	Value
Total crates indexed	245,551
Crates analyzed	53,544 (21.8%)
Dependency relationships	1,551,031
Total ecosystem downloads	266.56 billion
Temporal coverage	Nov 2014 – Mar 2026
Earliest crate	2014-11-11
Latest update	2026-03-30

Table 7.1: Dataset Summary Statistics

Percentile	Downloads
25th (P25)	1,149
50th (Median)	3,041
75th (P75)	10,667
95th (P95)	266,807
Mean	1,085,553

Table 7.2: Download Distribution Statistics

The download distribution is extremely skewed, a small number of crates are responsible for most downloads. This is probably due to popular dependencies and dependency chains.

We can see consistent growth in the number of crates, with a growth of 40% since 2015, and a significant growth in the last 2 years, especially since this data is from March.

Those crates represent are the most critical ones which would have the biggest impact in case of a compromise or vulnerability, affecting a very large number of dependent crates and their users.

Year	Crates Created
2014	598
2015	2,595
2016	3,596
2017	5,485
2018	7,716
2019	11,426
2020	17,142
2021	19,909
2022	25,582
2023	32,606
2024	36,126
2025	50,589
2026 (partial)	32,181

Table 7.3: Crates Created per Year

Crate	Dependents
serde	80,031
serde_json	55,479
thiserror	42,910
tokio	39,830
clap	33,224
anyhow	30,862
log	26,509
tracing	24,714
chrono	24,470
rand	19,968

Table 7.4: Most Depended-Upon Crates

7.2 Vulnerability Landscape

The cargo audit module identified a large number of known vulnerabilities across this small subset of the ecosystem.

A high proportion of vulnerabilities have no CVSS score.

The most common vulnerability is RUSTSEC-2026-0007 which is a memory corruption vulnerability in the "bytes" crate, which can lead to a buffer overflow and potential RCE.

Radicle packages suggest shared vulnerable dependencies relied upon.

7.3 Secret Leakage Findings

The gitleaks module detected substantial exposure across the analyzed repositories, a significant number is of false positives from test files.

The minimum of 0 bits is due to the system failing on select repositories.

'generic-api-key' detections are the most common with a percentage of 99.57%, indicating a high number of false positives from crypto constants and test files.

Metric	Value
Total vulnerability instances	19,558
Unique RustSec advisories	281
Affected crates	4,473 (8.4% of analyzed)
Crates with ≥ 1 vulnerability	8.36%
Average vulnerabilities per affected crate	4.37
Maximum vulnerabilities in single crate	62

Table 7.5: Vulnerability Statistics Summary

Severity Level	CVSS Score Range	Count	Percentage
Critical	90–100	415	2.1%
High	70–89	2,155	11.0%
Medium	40–69	1,583	8.1%
Low	1–39	352	1.8%
Unknown/Unscored	0	15,053	77.0%

Table 7.6: CVSS Severity Distribution

7.4 Build Script Risk Assessment

The build script analysis revealed the following:

Network calls, link directives and process spawning in FFI and sys crates is usually expected but the combination of these together with free TLDs is suspicious and should be manually investigated.

7.5 LLM Malicious Code Analysis

LLM analysis was disabled due to computational constraints on this test run

All crates received the default score of 0. The following runs will have a partial LLM analysis.

7.6 Typosquatting Detection Results

The majority of the high scores are legitimate forks or versions rather than typosquats. However pairs like 'windows-string' and 'windows-strings' are suspicious and should be manually investigated.

7.7 Dependency Chain Risk Propagation

Large dependency trees increase complexity of security audits and the attack surface, this list demonstrates potential supply chain risks, we can remove unnecessary dependencies if possible.

7.8 Performance and Scalability

The framework was deployed on a single machine with multiple workers.

RustSec ID	Occurrences
RUSTSEC-2026-0007	1,377
RUSTSEC-2026-0049	1,231
RUSTSEC-2026-0009	1,081
RUSTSEC-2024-0421	987
RUSTSEC-2025-0055	816
RUSTSEC-2026-0037	683
RUSTSEC-2023-0071	668
RUSTSEC-2025-0009	623
RUSTSEC-2024-0006	514
RUSTSEC-2026-0067	417

Table 7.7: Most Frequent RustSec Advisories

Crate	Vulnerability Count	Weighted Risk Score	Downloads
radicle-registry-node	58	1,796	1,819
radicle-registry-client	58	1,796	1,685
radicle-registry-runtime	58	1,796	3,119
memzero	36	1,631	—
keccak-hashier	36	1,631	—
tetcore-std	62	1,463	52,681
noble-contracts-proc-macro	62	1,463	7,372
tetcore-test-utils-test-crate	62	1,463	1,586

Table 7.8: High-Risk Crates by Vulnerability and Score

The configuration achieved a decently high throughput on the test dataset. The maximum of 33 minutes is indicative of the fact that in the ran version there was no cleaning of caches and the machine remained with no available storage, this was fixed and will not be present in the next runs.

Metric	Value
Total secrets detected	393,059
Unique detection rules triggered	31
Affected crates	1,496 (2.8% of analyzed)
Average secrets per affected crate	262.7

Table 7.9: Secret Detection (Gitleaks) Summary

Rule ID	Count	Percentage
generic-api-key	391,361	99.57%
aws-access-token	561	0.14%
curl-auth-header	426	0.11%
private-key	244	0.06%
jwt	109	0.03%
github-pat	73	0.02%
curl-auth-user	57	0.01%
stripe-access-token	48	0.01%
datadog-access-token	32	0.01%
slack-webhook-url	31	0.01%

Table 7.10: Gitleaks Rule Frequency Distribution

Metric	Value
Average entropy	3.82 bits
Minimum entropy	0.00 bits
Maximum entropy	5.68 bits

Table 7.11: Entropy Statistics for Detected Secrets

Crate	Secret Count
xoofff	187,640
fuel-core-bft	29,431
libjade-sys	9,646
libcrux-hacl	9,646
libcrux-platform	9,646
tnt-core-fixtures	8,960
cryptocol	4,943
tinyaes	4,372
tari_base_node_client	2,892
minotari_wallet ffi	2,892

Table 7.12: Crates with Highest Number of Detected Secrets

Indicator	Count	Percentage of Analyzed
Total crates with build.rs	2,463	4.6%
Network calls (reqwest, curl, etc.)	246	0.46%
Link directives	1,084	2.02%
Process spawning (Command::new)	730	1.36%
Raw IP addresses	0	0.00%
Free/suspicious TLDs	13	0.02%
Crates with executable files	841	1.57%

Table 7.13: Build Script Risk Indicators

Entropy Range	Count
Very High (≥ 5.5)	14
High (5.0–5.5)	380
Medium (4.5–5.0)	1,869
Low (0–4.5)	200
None (no build.rs)	51,081

Table 7.14: Build Script Entropy Distribution

Crate	Entropy Score
shiguredo_fdk_aac	5.91
shiguredo_dav1d	5.84
lbl-tg-rcore-tutorial-vga	5.80
lbl-tg-rcore-tutorial-ch1-game	5.80
tg-rcore-tutorial-ch1-t2l3	5.80
shiguredo_nvcodec	5.79
shiguredo_webrtc	5.77
shiguredo_opus	5.72
shiguredo_libyuv	5.71
shiguredo_aom	5.70

Table 7.15: Highest Entropy Build Scripts

Crate	Network	Link	Process	Free TLDs	Total Indicators
libsodium-ffi	✓	✓	✓	—	3
ghidra-decompiler-sys	✓	✓	✓	—	3
tcmalloc2	✓	✓	✓	—	3
intel-onemkl-prebuild	✓	✓	✓	—	3
jc-libavif-sys	✓	✓	✓	—	3
libntirpc-sys	✓	✓	✓	—	3
libtor-src	—	✓	✓	✓	3

Table 7.16: Crates with Multiple Build.rs Risk Indicators (≥ 3)

Score Range	Count	Percentage
Minimal (0–24)	53,544	100.0%
Low (25–49)	0	0.0%
Medium (50–74)	0	0.0%
High (75–100)	0	0.0%

Table 7.17: LLM Score Distribution

Metric	Value
Total similar pairs identified	37,596
Average combined score	79.16
Score range	75–98

Table 7.18: Typosquatting Detection Overview

Score Range	Count	Percentage
Very High (90–100)	882	2.3%
High (80–89)	13,286	35.3%
Medium (75–79)	23,428	62.3%

Table 7.19: Typosquatting Score Distribution

Crate	Similar To	Score
opentelemetry-appender-tracing2	opentelemetry-appender-tracing	98
tauri-plugin-notifications	tauri-plugin-notification	97
derive_destructure	derive_destructure2	97
dioxus-config-macros	dioxus-config-macro	97
safer_ffi-proc_macro	safer_ffi-proc_macros	97
dprint-plugin-jsonc	dprint-plugin-json	97
clap-verbosity-flag2	clap-verbosity-flag	97
protobuf-codegen-pure3	protobuf-codegen-pure	97
proc-macro-error-attr2	proc-macro-error-attr	97
tauri-plugin-shellx	tauri-plugin-shell	97

Table 7.20: Highest Scoring Similar Typosquatting Pairs

Potential Typosquat	Target Crate	Typosquat DL	Target DL	Score
opentelemetry-appender-tracing2	opentelemetry-appender-tracing	5,043	6.4M	98
proc-macro-error-attr2	proc-macro-error-attr	75.8M	258.9M	97
azure_storage_blob	azure_storage_blobs	735K	16.1M	97
clap-verbosity-flag2	clap-verbosity-flag	3,490	12.0M	97
windows-string	windows-strings	1,670	148.8M	96
async-recursion2	async-recursion	2,475	95.7M	96
tikv-jemalloc-sys2	tikv-jemalloc-sys	1,639	66.8M	96

Table 7.21: High-Risk Typosquats Targeting Popular Crates (1M+ Downloads)

Metric	Value
Crates with malicious dependencies flag	0
Crates with executable files	841 (1.57%)
Average dependencies per crate	6.32

Table 7.22: Transitive Risk Summary

Crate	Direct Dependencies
pezkuwi-sdk	394
polkadot-sdk	380
zigbee2mqtt-types	254
gnostr	199
forest-filecoin	169
apollo-router	150
uhg-custom-apollo-roouter	143
solana-core	141
zero4rs	140
solana-runtime	136

Table 7.23: Crates with Most Dependencies

Metric	Value
Total crates analyzed	52,550
Total wall-clock time	14h 24m 5s
Total CPU time	60.09 hours
Throughput	1.01 crates/second
Throughput	60.8 crates/minute
Throughput	3,649 crates/hour

Table 7.24: Analysis Performance Summary

Module	Average Duration (ms)
Total per crate	4,116.58
Cargo audit	5,540.03
Gitleaks	1,047.31
Build.rs analysis	0.03
LLM analysis	0.00 (disabled)
Git clone/download	1,645.94

Table 7.25: Per-Analysis Module Timing

Percentile	Duration (ms)
P50 (Median)	1,077
P95	12,056
P99	88,183
Maximum	1,972,283 (32.9 min)

Table 7.26: Analysis Duration Percentiles

Worker ID	Crates Analyzed	Avg Duration (ms)
admin-pc	52,550	4,116.58

Table 7.27: Worker Performance

Chapter 8

Discussion

8.1 Interpretation of Results

The key findings from the analysis is the following:

8.4% of analyzed crates contain at least one vulnerability, with an average of 4.37 per vulnerable crate. This is lower than the npm rate, however it remains a concern

The top 10 dependent crates have over 350.000 dependent crates.

Over 393.059 secrets were detected by gitleaks, however the majority of them are false positives of generic API key patterns, and probably represent test and constants. The non-generic detections are a small number and should require manual analysis.

Around 4.6% of crates contain build scripts, with over 240 performing network operations.

8.2 Comparison with Existing Approaches

TODO: After rerun of framework

8.3 Limitations

8.3.1 Detection Limitations

False negatives:

- Novel patterns and attacks can evade our heuristics
- Custom encoding might be done to try to evade obfuscation detection
- Social engineering and account overtaking are not addressed

False positives:

- Legitimate crates with unusual behavior that still remain benign
- Test code may trigger alerts
- High entropy with legitimate usecases
- Common naming patterns can produce false alerts of the typosquatting features

8.3.2 LLM-Specific Limitations

- LLMs might hallucinate vulnerabilities or hallucinate that benign code is malicious or viceversa

- The very limited context window is a problem due to some libraries having large codebases, which are chunked on the current implementation, which can make the model miss cross-file malicious behavior (behavior that present in two separate files is benign on it's own but either imports or api usage makes it potentially malicious)

- Reproducibility is another limitation since LLM outputs vary even with low or zero temperature parameter.

- Prompt injection is another problem, adversaries might contain specially encoded unicode characters or try to do prompt injection in comments to evade the model

8.3.3 Scalability and Resource Limitations

LLm inference is extremely computationally expensive for an ecosystem deployment.

Analysis time can take a very long time, especially on large crates, adversaries could do a DoW attack by spamming fake crates with large codebases to attack crates.io.

Storage is another problem if storing crates is required, besides us only cloning with a depth of one and only one branch, some repositories are still large

Besides storage, network bandwidth is also a potential problem, especially if the git services used rate limit us.

8.3.4 Ground Truth Limitations

Validation is extremely hard to do, since only a very small number of existing crates have been documented in a public way as malicious

We also have an unknown rate for false negatives since we cannot know how many attacks exists undetected

8.3.5 Generalizability Limitations

The framework is designed for crates.io and Rust, adoption for other ecosystems would require a full rewrite.

Besides the tools like cargo audit being adapted for Rust, also the build.rs heuristics are also language specific

8.4 Threats to Validity

The CVSS scoring and string similarity algorithms have been tested via unit tests, however we can never know if there are any edge cases

Results are limited to crates with valid repository links in the metadata, some crates have been moved from the link from the metadata and moved to docs.rs (i.e. Find_Matt_C)

The LLM score is only a guide and cannot guarantee anything, be it either the presence of malivious behavior or its absence.

Chapter 9

Conclusions and Future Work

9.1 Summary of Contributions

The thesis makes the following significant contributions:

- Comprehensive analysis framework - Implemented multiple analysis modules each with a focus on different techniques
- Distributed architecture - The redis work queues and topological ordering, together with the producer-consumer design and parallel processing we enable horizontal scaling achieving high throughput
- Analysis modules - Composite typosquatting scoring, heuristic analysis of build scripts targeting language specific attack vectors
- Empirical security assesment - Provide useful insights into the security landscape of the entire ecosystem.
- Web dashboard - Interactive and friendly exploration of results

9.2 Future Directions

We discovered several possible future directions:

- Analysis of specific crates versions, tracking vulnerabilities of a crate based on the specific version of the dependent.
- Real-time monitoring can be improved, continuous scanning and alerting on new findings.
- Improved LLM integration for usage of advanced providers, better prompting and usage of more advanced models.
- Sandboxed build script execution for hybrid analysis, implementing a sandboxed environment would help to execute the script in a safe manner and extract information about the behavior providing better classification
- Collaborative review would be great to integrate with cargo crev to notify maintainers or contributors about possible security issues
- Caching of repositories would greatly improve the performance of the system, by avoiding the need to clone each repository every time.

Appendix A

Database Schema

```
CREATE TABLE crates (  
    id BIGINT PRIMARY KEY,  
    name TEXT NOT NULL,  
    repository TEXT NOT NULL,  
    crate_downloads BIGINT NOT NULL,  
    db_created_at TIMESTAMP NOT NULL DEFAULT CURRENT_TIMESTAMP,  
    db_updated_at TIMESTAMP NOT NULL DEFAULT CURRENT_TIMESTAMP,  
    crate_created_at TIMESTAMP NOT NULL,  
    crate_updated_at TIMESTAMP NOT NULL  
);
```

```
CREATE TABLE dependencies (  
    crate_id BIGINT NOT NULL REFERENCES crates(id),  
    dependency_id BIGINT NOT NULL REFERENCES crates(id),  
    db_created_at TIMESTAMP NOT NULL DEFAULT CURRENT_TIMESTAMP,  
    db_updated_at TIMESTAMP NOT NULL DEFAULT CURRENT_TIMESTAMP,  
    PRIMARY KEY (crate_id, dependency_id)  
);
```

```
CREATE TABLE scan_results (  
    id BIGINT PRIMARY KEY REFERENCES crates(id),  
    has_malicious_dependencies BOOLEAN NOT NULL,  
    llm_malicious_score SMALLINT NOT NULL,  
    llm_notes TEXT,  
    has_executable_files BOOLEAN,  
    cargo_audit_max_dep_score SMALLINT,  
    cargo_audit_vulns_count SMALLINT,  
    build_rs_network_calls BOOLEAN,  
    build_rs_has_link_directive BOOLEAN,  
    build_rs_entropy_score DOUBLE PRECISION,  
    build_rs_has_process_spawning BOOLEAN,  
    build_rs_has_raw_ip BOOLEAN,  
    build_rs_has_free_tlds BOOLEAN,  
    entropy_score DOUBLE PRECISION  
);
```

```
CREATE TABLE cargo_audit_results (  
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
    crate_ BIGINT REFERENCES crates(id),  
    rustsec_id VARCHAR(10) NOT NULL,  
    severity SMALLINT  
);
```

```
CREATE TABLE gitleaks_results (  
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
    crate_ BIGINT REFERENCES crates(id),  
    rule_id VARCHAR(255) NOT NULL,  
    secret VARCHAR(255) NOT NULL,  
    loc VARCHAR(255) NOT NULL,  
    entropy DOUBLE PRECISION NOT NULL  
);
```

```
CREATE TABLE typosquat_results (  
    id BIGSERIAL PRIMARY KEY,  
    crate_id BIGINT NOT NULL REFERENCES crates(id),  
    similar_crate_id BIGINT NOT NULL REFERENCES crates(id),  
    levenshtein_score SMALLINT NOT NULL,  
    damerau_levenshtein_score SMALLINT NOT NULL,  
    jaro_winkler_score SMALLINT NOT NULL,  
    keyboard_distance_score SMALLINT NOT NULL,  
    prefix_similarity_score SMALLINT NOT NULL,  
    combined_score SMALLINT NOT NULL,  
    db_created_at TIMESTAMP NOT NULL DEFAULT CURRENT_TIMESTAMP,  
    UNIQUE (crate_id, similar_crate_id),  
    CHECK (crate_id != similar_crate_id)  
);
```

```
CREATE TABLE runner_metadata (  
    run_time TIMESTAMP PRIMARY KEY,  
    last_checked_crate BIGINT REFERENCES crates(id)  
);
```

Appendix B

Configuration Reference

```
# Database
max_connections = 5
connection_string = "postgres://postgres:postgres@localhost:5432/crates"

# Redis
redis_url = "redis://localhost:6379"

# Data Import
update_source_data = true
use_only_first_x_crates = 0          # 0 = all crates
max_crates_downloads = 0            # 0 = no filter
insert_crates_in_chunks = 4000

# LLM Configuration
llm_enabled = false
lm_studio_url = "http://localhost:1234/v1"
lm_studio_model = "qwen/qwen2.5-coder-14b"
llm_max_tokens = 4096
llm_temperature = 0.1
llm_max_context_chars = 24576
llm_malicious_score_threshold = 75
```


Appendix C

Web Dashboard API Reference

TODO: Make webpanel use sailfish

Endpoint	Method	Description
/api/stats	GET	Dashboard aggregate statistics
/api/stats/severity	GET	Vulnerability severity distribution
/api/stats/risk	GET	LLM risk score distribution
/api/stats/build-rs	GET	Build script indicator statistics
/api/stats/top-downloaded	GET	Top 10 crates by download count
/api/stats/gitleaks-rules	GET	Top 10 Gitleaks rules by frequency
/api/crates	GET	Paginated crate list (filterable)
/api/crates/{id}	GET	Single crate by ID
/api/crates/name/{name}	GET	Single crate by name
/api/dependencies	GET	Paginated dependency list
/api/scan-results	GET	Paginated scan results (filterable)
/api/scan-results/{id}	GET	Single scan result
/api/cargo-audit	GET	Paginated audit results
/api/gitleaks	GET	Paginated secret findings
/api/typosquat	GET	Paginated typosquatting pairs

Table C.1: API Endpoints

Appendix D

Artifact Documentation and Reproducibility

D.1 Software Artifacts

Artifact	Description	Location
crates_downloader	Main analysis framework (Rust)	crates_downloader/
webpanel	Dashboard web application (Rust + Axum)	webpanel/
Database migrations	PostgreSQL schema definitions	migrations/
Docker configuration	Container orchestration	docker-compose.*.yaml, D
Configuration template	Analysis parameters	config.toml

D.2 Data Artifacts

Artifact	Description	Format
crates.io dump	Input data from crates.io	CSV files
Analysis results	Complete ecosystem scan	PostgreSQL database dump
Vulnerability findings	cargo audit results	JSON / PostgreSQL
Secret detections	Gitleaks findings	JSON / PostgreSQL
Typosquatting pairs	Similarity computations	PostgreSQL

D.3 Hardware Requirements

Component	Minimum	Recommended
CPU	4 cores	16+ cores
RAM	16 GB	64 GB
Storage	500 GB SSD	1 TB NVMe + tmpfs
GPU	None (CPU inference)	NVIDIA RTX 5000 series for fast inference
Network	100 Mbps	1 Gbps

D.4 Software Requirements

Software	Version	Purpose
Rust	(stable)	Core application
PostgreSQL	17+	Data storage
Redis	7+	Work queue
Docker	24+	Containerization
Docker Compose	2.20+	Orchestration
Gitleaks	8+	Secret detection
LM Studio	0.2.20+	Local LLM hosting
Git	2.40+	Repository cloning

D.5 Reproduction Steps

1. Clone the repository:

```
git clone https://github.com/[REDACTED]/crates-security-framework.git
cd crates-security-framework
```

2. Configure environment:

```
cp config.toml.example config.toml
# Edit config.toml with database and Redis URLs
```

3. Start infrastructure:

```
docker-compose up -d postgres redis
```

4. Initialize database:

```
diesel migration run
```

5. Download crates.io database dump:

```
cargo run --release -- update-database
```

6. Build dependency graph:

```
cargo run --release -- run-analysis
```

7. Start workers (adjust count as needed):

```
docker-compose -f docker-compose.workers.yml up -d --scale worker=10
```

8. Monitor progress:

```
cargo run --release -- queue-status
```

9. **Run post-processing:**

```
cargo run --release -- compute-dependencies
```

10. **Run typosquatting analysis:**

```
cargo run --release -- typosquat --min-score 75 --top-crates 10000
```

11. **Start web dashboard:**

```
cd webpanel && cargo run --release
```

D.6 Expected Outputs

PostgreSQL database containing:

- All crate metadata from crates.io
- Complete dependency graph
- Scan results for all analyzed crates
- Individual vulnerability findings (cargo_audit_results)
- Individual secret findings (gitleaks_results)
- Typosquatting pair candidates (typosquat_results)

Log files in logs/

Working web dashboard

Appendix E

Glossary

Term	Definition
C2 (Command and Control)	Server infrastructure used by attackers to communicate with compromised systems
CVSS	Common Vulnerability Scoring System, a standardized severity rating for vulnerabilities
DAG	Directed Acyclic Graph, used to model dependency relationships
Entropy	Shannon entropy, a measure of randomness/information density in data
Gitleaks	Open-source tool for detecting secrets in Git repositories
LLM	Large Language Model, an AI model trained on large text corpora
RustSec	The Rust security advisory database
Transitive dependency	Dependencies of dependencies (indirect dependencies)
Typosquatting	Registering names similar to popular packages to deceive users
YARA	Pattern matching tool commonly used for malware detection

Table E.1: Glossary of Terms

Appendix F

Generative AI Tools Usage

An overview how generative AI Tools are used in this thesis is presented in Table ??.

AI Tool	Use case	Used in	Remarks
ChatGPT 5.0	Provide remarks of writing, reviewing typos	All chapters	Prompts in section ??
	Generate initial state of the art bibliography list	Chapter 2	Prompts in section ??
Microsoft Copilot	Code generation	Chapter 4 pp.2-21	Details in section ??
Deep L	Translation of text pages	Chapter 2	Prompts in section ??
Grammarly	Correct grammar and style mistakes	Entire work	Free version
...			

Table F.1: Overview of generative AI Usage

F.1 ChatGPT Details

Specify for each usage the following:

1. Purpose of use
2. Method (which prompts and for which parts of the thesis)
3. How the output has been examined/verified (has the generated output been examined for accuracy and reliability?)
4. How the output is implemented/used

F.2 Deep L Details**F.3 Microsoft Copilot**